

The Living Craft

Sunil content rewrite review

STATUS	OWNER	LAST UPDATED
Prepared for review	Team Alchemy	10 September 2026

Authors	Team Alchemy
Reviewers	Sunil Mathew and Ein
Related docs	Cohort handoff package
Scope	Fifteen originals and proposed alternatives with review notes

Prepared by Team Alchemy for Sunil Mathew. 10 September 2026. All fifteen originals remain the active baseline. These are proposed alternatives preserving each topic, audience, format, sequence and destination. The company's fifteen posts are unchanged.

REVIEW SEQUENCE

Begin with D01 written, D05 carousel and D03 video. These three samples were prepared first to establish voice; the remaining twelve extend that direction and do not imply sample approval. Review narrative delivery and technical accuracy, then select versions by ID.

VOICE AND EVIDENCE

An experienced mentor and hands-on builder, with conversational explanations and light situational humour. Scenarios are illustrative unless a source establishes an actual event. Existing calls support the programme reasoning, but do not establish personal workplace anecdotes. No new biographical stories are claimed.

LC-OCT-D01 The boundary your demo leaves open

Day 1 / Written / Member / Original active

Current version	Proposed rewrite	Review note
An agent can recommend a refund without being allowed to issue it. The moment you connect the payment tool, the design question changes. I would ask which evidence permits the action, how repeated requests are handled and who can intervene when the answer is uncertain. Those decisions belong in the architecture, before a convincing demo becomes a team commitment. At The Living Craft, members build a working agentic system and practise explaining the choices behind it. The aim is to strengthen the ability to design, critique and guide a team's architecture decisions. If you already work with system design and want to develop that responsibility, explore the open cohort. Programme and application: https://learning.thelivingcraft.ai/?utm_source=linkedin&utm_medium=organic_social&utm_campaign=cohort_enterprise_v3&utm_content=lc-oct-d01	Imagine a returns assistant giving a perfectly sensible answer: the customer should get a refund. Everyone can follow the reasoning. Then someone connects the payment tool. Now the assistant can move money. That small addition deserves a much bigger conversation. Who allowed the refund? What evidence was enough? If the customer asks twice, do we pay twice? When the answer is uncertain, who gets the next decision? A good demo makes the happy path easy to see. I want the design to make these less comfortable questions visible too. That is the connection we practise at The Living Craft: build a working agentic system, examine what it does, and explain the choices behind it. If you already work with system design and want to get better at making and guiding those decisions, explore the open cohort. Programme and application: https://learning.thelivingcraft.ai/?utm_source=linkedin&utm_medium=organic_social&utm_campaign=cohort_enterprise_v3&utm_content=lc-oct-d01	Illustrative returns scenario, not a claimed customer incident. Replaces an abstract opening with a moment when the system gains authority. Sunil to confirm natural phrasing and technical explanation. Pending review. Original active.

LC-OCT-D03 Where should an agent stop?

Day 3 / Video / Member / Original active

Current version	Proposed rewrite	Review note
Before I ask which model an agent uses, I want to understand its authority. In this illustrative returns workflow, recommending a refund and issuing one are different responsibilities. A boundary makes that difference visible to the team and testable in the system. The Living Craft brings this kind of reasoning into a working build, evaluation and design review. Explore the open cohort if these are decisions you want to own. Programme and application: https://learning.thelivingcraft.ai/?utm_source=linkedin&utm_medium=organic_social&utm_campaign=cohort_enterprise_v3&utm_content=lc-oct-d03	A refund recommendation and a refund payment can look like one step on a whiteboard. The customer's bank account would disagree. In this illustrative example, I walk through what changes when an agent gains permission to act. The interesting questions are about evidence, limits and who handles uncertainty. We practise this reasoning through a working build at The Living Craft. Explore the open cohort if these are decisions you want to own. Programme and application: https://learning.thelivingcraft.ai/?utm_source=linkedin&utm_medium=organic_social&utm_campaign=cohort_enterprise_v3&utm_content=lc-oct-d03	Illustrative scenario with light situational humour; no personal incident asserted. Preserves authority and action-boundary teaching plus cohort invitation. Sunil to check the humour and delivery aloud. Pending review. Original active.
Spoken script Before I ask which model an agent uses, I want to understand what it is allowed to do. Imagine a retail returns assistant. It reads a policy, looks at the request and recommends a refund. Now connect a tool that can issue that refund. The task may sound similar, but the authority has changed. I would separate the recommendation from the action. What evidence permits the refund? Which limits must the tool enforce? Where does an uncertain request go? And what happens if the same request arrives twice?	Spoken script Imagine a returns assistant. It reads the policy, checks a request and recommends a refund. So far, you can look at its answer and decide what to do. Then we connect a tool that can issue the refund. It looks like one more box on the diagram. But that box can move money. That is where I would stop and ask a few questions. What evidence permits the action? What limits does the tool enforce? If the same request arrives again, can we recognise it? And if the policy does not give us a clear answer, who takes over? The recommendation and the payment need different controls. Writing a	Same format and cohort purpose.

Current version	Proposed rewrite	Review note
These are architecture questions the team should be able to explain. A prompt can describe the intended behaviour, but the action boundary also needs to exist in the system. This is an illustrative example of the work we discuss at The Living Craft. Members build a working agentic system, examine its behaviour and practise defending the design choices behind it. If you already work with system design and want to develop that judgment, explore the open cohort.	reassuring instruction in the prompt does not, by itself, create those controls in the system. This is an illustrative example, but it gives us a concrete design to discuss. We can build it, examine its behaviour and change our decisions when the evidence calls for it. That is the kind of practice we bring into The Living Craft. If you already work with system design and want to strengthen how you design, critique and guide agentic systems, explore the open cohort.	

LC-OCT-D05 Fit is about responsibility

Day 5 / Carousel / Member / Original active

Current version	Proposed rewrite	Review note
A title alone does not tell me whether this programme is the right fit. Can you explain a system you have worked on? Can you name a constraint that changed its design? Are you ready to build, examine evidence and revise a decision? Prior system-design exposure matters. You do not need to arrive with a company problem ready for consulting. The programme uses industry-relevant practice to develop your own design judgment. Describe your experience and learning goal in the open-cohort application. Programme and application: https://learning.thelivingcraft.ai/?utm_source=linkedin&utm_medium=organic_social&utm_campaign=cohort_enterprise_v3&utm_content=lc-oct-d05	A job title tells me very little about the design conversation we could have. A trade-off tells me much more. For example, imagine describing a system that became easier to operate after the team gave up a feature. What made that choice sensible? What did it cost? Would you make it again? Now we have something useful to discuss. That is the starting point I care about for The Living Craft. Prior system-design exposure, a question worth working on, and willingness to build and revise through feedback. You can work with an industry case. You do not need to bring a confidential company project. If that sounds like your next step, describe your experience and learning goal in the open-cohort application. Programme and application: https://learning.thelivingcraft.ai/?utm_source=linkedin&utm_medium=organic_social&utm_campaign=cohort_enterprise_v3&utm_content=lc-oct-d05	Illustrative trade-off, not a biographical anecdote. Six panels retain the original fit, experience, practice and application purpose. Sunil to confirm that the example sounds like his explanation of readiness. Pending review. Original active.
Creative copy Is this your next step? Fit starts with the decisions you already understand. Explain a system Describe its purpose, components and constraints. Name a trade-off	Creative copy Start with a decision Your title starts the introduction. A design choice starts the conversation. Explain a system What did it need to do, and which constraint shaped it? Tell us about a trade-off	Same format and cohort purpose.

Current version	Proposed rewrite	Review note
What did you choose, and what did that choice cost? Make room for practice Live discussion and independent work both matter. Bring a learning goal You can work on an industry case. A company project is not a prerequisite. Apply with context Describe your design experience at learning.thelivingcraft.ai	What became better? What did you give up? Be ready to revise Building and feedback can change the answer you started with. Bring a learning goal An industry case is enough. Keep confidential company material out. Start the fit conversation Describe your experience and goal at learning.thelivingcraft.ai	

LC-OCT-D07 A course choice can be a funding decision

Day 7 / Written / Member / Original active

Current version	Proposed rewrite	Review note
You may be ready to join a programme and still need someone else to approve the time and funding. I would make that request specific: which architecture responsibility are you taking on, what do you need to practise and how will your team use your stronger judgment? The Living Craft combines a working agentic system, design reasoning and continuous feedback. The commitment is 30 live hours plus independent work, with the detailed schedule and workload confirmed before joining. Use the manager summary to start the employer-funding conversation. It should help your manager assess a real responsibility, not rely on a certificate or a promised promotion. For the employer-funding summary, email apply@thelivingcraft.ai and mention LC-OCT-D07.	You find a programme that looks right. Then comes the second application: the one to your manager. The questions change. What work will this help with? Where will the time come from? Who approves the funding? I would start with one responsibility you already hold. Perhaps you need to review the team's agent design and you want more practice connecting a design choice to evidence. That gives the funding conversation something concrete. The Living Craft combines a working agentic system, architecture reasoning and continuous feedback. It includes 30 live hours plus independent work; confirm the detailed schedule and workload before making the commitment. The manager summary can help you explain that connection. Use it to begin the employer-funding conversation for an open-cohort place. For the employer-funding summary, email apply@thelivingcraft.ai and mention LC-OCT-D07.	Uses a recognisable funding situation, not an invented personal incident. Keeps the manager-summary route and full commitment. Sunil to confirm the conversational opening. Pending review. Original active.

LC-OCT-D09 A useful evaluation changes a decision

Day 9 / Carousel / Member / Original active

Current version	Proposed rewrite	Review note
I want an evaluation to change something the team can name. Consider an illustrative assistant that handles returns. A polite final answer is not enough if it used the wrong policy or repeated a refund action. Review the task, the evidence, the tool behaviour and the outcome separately. Then decide what to change: a permission, a workflow, an evaluator or a human review point. The value is in that connection between evidence and architecture. Explore the open cohort to practise this with a working system. Programme and application: https://learning.thelivingcraft.ai/?utm_source=linkedin&utm_medium=organic_social&utm_campaign=cohort_enterprise_v3&utm_content=lc-oct-d09	Imagine an assistant ending a returns conversation with a lovely answer. Clear, polite, reassuring. Then you discover it used the wrong policy. The customer liked the conversation. The review now has a different problem to solve. That is why I want to look behind the final answer. What was the task? Which evidence did the system use? What did the tool actually do? Was the outcome right? Once we can see that path, we can make a useful change: narrow a permission, fix the workflow or add a review point. At The Living Craft, we practise that connection between evaluation and architecture through a working system. Explore the open cohort if you want to get better at making those calls. Programme and application: https://learning.thelivingcraft.ai/?utm_source=linkedin&utm_medium=organic_social&utm_campaign=cohort_enterprise_v3&utm_content=lc-oct-d09	Illustrative policy failure; no client result asserted. Keeps evaluation-to-design purpose and six-panel structure. Sunil to validate the example and action language. Pending review. Original active.
Creative copy What would change your design? An evaluation is useful when it informs a decision. Start with the task Illustrative example: resolve a returns request under the right policy. Inspect the path	Creative copy A lovely answer can still be wrong Imagine a returns assistant using the wrong policy. Start with the task What should a correct outcome look like for this request? Follow the evidence	Same format and cohort purpose.

Current version	Proposed rewrite	Review note
What evidence and tools did the agent actually use? Check the consequence A plausible answer can hide an incorrect or repeated action. Name the next change Adjust a boundary, workflow or review point, then evaluate again. Practise the connection Build and critique an agentic system. learning.thelivingcraft.ai	Which policy did the assistant find and use? Check the action What did the tool do? Could a repeated request repeat the refund? Make a decision Change the boundary, workflow or review point. Then examine the result. Practise with a working system Explore agentic-system design and critique at learning.thelivingcraft.ai	

LC-OCT-D11 What evidence would you ask for?

Day 11 / Video / Member / Original active

Current version	Proposed rewrite	Review note
A release discussion needs evidence about behaviour, not only an impressive example. I would want to see where the system succeeds, what it fails on and what happens when its tools or evidence are uncertain. A team should be able to explain which result would change its decision. In The Living Craft, evaluation supports the practical build and architecture review. Explore the open cohort to develop that judgment. Programme and application: https://learning.thelivingcraft.ai/?utm_source=linkedin&utm_medium=organic_social&utm_campaign=cohort_enterprise_v3&utm_content=lc-oct-d11	A demo is usually a short visit. Your system has to live there. Imagine the useful next question being: what happened on the requests we did not choose for the presentation? I want to see failures, uncertain evidence and awkward tool results, as well as the examples that worked. Then we can discuss which finding would change the design. That is how evaluation supports the practical build at The Living Craft. Explore the open cohort to develop that judgment. Programme and application: https://learning.thelivingcraft.ai/?utm_source=linkedin&utm_medium=organic_social&utm_campaign=cohort_enterprise_v3&utm_content=lc-oct-d11	Illustrative demo review, not a story about a past team. Adds a memorable opening while preserving evidence and cohort purpose. Sunil to check delivery and technical interpretation. Pending review. Original active.
Spoken script What evidence would make you comfortable with an agent's next release? I would start with the task it is supposed to perform, then look at the behaviour behind the final answer. Did it use the right evidence? Did it call the permitted tools? What happened when a tool failed or returned an uncertain result? For an illustrative returns assistant, a polite answer is not enough. I would want to know whether it applied the relevant policy and whether a repeated request could lead to a repeated action. Then I would ask the team which finding would change the design. Maybe the workflow needs a review point. Maybe	Spoken script Imagine a team showing an agent that handles a returns request beautifully. The answer is clear. The tool works. Everyone can see the idea. Now I would ask to see the requests we did not put in the demo. What happened when the policy was unclear? What happened when the tool failed? Did the assistant recognise a repeated request, or was it about to repeat the action? Those questions are not there to spoil a good demonstration. They help us understand what we have actually learned from it.	Same format and cohort purpose.

Current version	Proposed rewrite	Review note
the permission is too broad. Maybe the evaluation is missing a class of failures. The useful part is the connection between evidence and an engineering decision. At The Living Craft, members practise that connection through a working agentic system, evaluation and design review. Explore the open cohort if this is the kind of responsibility you want to strengthen.	I would want the team to connect each finding to a decision. If this class of requests keeps failing, what do we change? The workflow? A permission? The evidence we retrieve? Or the evaluation itself? Sometimes the next step is a small test, not a bigger design. But we should be able to explain why that test matters. This returns example is illustrative. At The Living Craft, members build a working agentic system and use evaluation evidence in architecture discussions. If you want more practice making and explaining those decisions, explore the open cohort.	

LC-OCT-D13 A cheaper response can cost more

Day 13 / Written / Member / Original active

Current version	Proposed rewrite	Review note
A lower model bill does not settle the cost of an agentic workflow. I would also examine retries, human review, latency and the consequence of an incorrect action. An apparently cheap path may create expensive rework. A more capable model may still be unnecessary for a bounded task. The design discussion needs a defined task and evidence. Decide which cost you are optimising and which risk you are willing to carry, then make the trade-off explicit. That is part of the judgment members practise while building and reviewing systems in The Living Craft. Explore the open cohort if these choices belong in your role. Programme and application: https://learning.thelivingcraft.ai/?utm_source=linkedin&utm_medium=organic_social&utm_campaign=cohort_enterprise_v3&utm_content=lc-oct-d13	Imagine celebrating a smaller model bill, then asking three people to fix what the workflow got wrong. The saving looked better before the rest of the cost arrived. That is why I would look beyond the price of one response. Retries, review time, delay and the consequence of an incorrect action all belong in the conversation. A more capable model is not automatically the answer either. A bounded task may work well with a simpler approach. First define the task. Then compare the behaviour and the full cost of getting the job done. Otherwise, we can optimise the easiest number to see and miss the work happening around it. We practise these trade-offs while building and reviewing agentic systems at The Living Craft. Explore the open cohort if these are choices you need to make. Programme and application: https://learning.thelivingcraft.ai/?utm_source=linkedin&utm_medium=organic_social&utm_campaign=cohort_enterprise_v3&utm_content=lc-oct-d13	Hypothetical rework scenario; three people is illustrative, not a measured result. No ROI claim. Sunil to confirm the cost explanation and whether he prefers a different everyday example. Pending review. Original active.

LC-OCT-D15 Feedback should tell you what to do next

Day 15 / Carousel / Member / Original active

Current version	Proposed rewrite	Review note
Useful feedback names a strength, a gap and a next piece of work. If a design handles a normal request but cannot explain a failed tool action, the follow-up should be specific. Add the missing behaviour, show the evidence and revisit the decision. The Living Craft uses continuous feedback across the work. The programme does not offer a certificate. The focus is the capability you develop through building, evaluating and reviewing a system. Explore the open cohort and discuss the experience with Sunil. Programme and application: https://learning.thelivingcraft.ai/?utm_source=linkedin&utm_medium=organic_social&utm_campaign=cohort_enterprise_v3&utm_content=lc-oct-d15	“Needs improvement” can keep someone busy for a very long time. It does not tell them where to start. Imagine a design that handles the normal request well but loses track of what happened when a tool fails. Now the feedback can be useful: keep the part that works, make the uncertain result visible, decide who handles it, and come back with the revised behaviour. The next conversation has something new to examine. That is what I want feedback to do at The Living Craft. Help you see a strength, locate a gap and work on a concrete next step. We keep returning to the working system and the reasoning behind it. The programme does not offer a certificate. Explore the open cohort if this kind of practice and feedback is what you need next. Programme and application: https://learning.thelivingcraft.ai/?utm_source=linkedin&utm_medium=organic_social&utm_campaign=cohort_enterprise_v3&utm_content=lc-oct-d15	Generic quoted feedback, not attributed to an actual person. Illustrative tool failure. Retains no-certificate fact and continuous-feedback purpose. Pending review. Original active.
Creative copy What should you work on next? Continuous feedback connects progress to action. Recognise a strength Which design decision can you explain with evidence? Locate the gap	Creative copy What do I change next? Good feedback gives you somewhere useful to start. Keep what works Name a decision you can explain with evidence. Find the loose end	Same format and cohort purpose.

Current version	Proposed rewrite	Review note
Which behaviour or assumption remains unresolved?	Illustrative example: the tool failed, but the system cannot tell what happened.	
Define follow-up work	Make the next step concrete	
Make the next change concrete enough to examine.	Expose the uncertain result and define the next decision.	
Return to the design	Bring the revision back	
Use the revised work to deepen the discussion. No certificate is offered.	Examine what changed. Continuous feedback; no certificate.	
Explore the experience	Explore the learning experience	
Programme fit and applications: learning.thelivingcraft.ai	Build, evaluate and revise at learning.thelivingcraft.ai	

LC-OCT-D17 Give a manager a decision they can make

Day 17 / Written / Member / Original active

Current version	Proposed rewrite	Review note
A funding request becomes easier to assess when it names the work. I would include the architecture responsibility you hold, one capability you want to strengthen, the time you need and how you intend to apply the learning. Confirm the programme schedule and independent-work expectation before asking for final approval. The Living Craft provides 30 live hours with practical building and continuous feedback. The manager summary explains the individual funding route separately from an enterprise group purchase. Use that summary to discuss employer funding for your open-cohort application. For the employer-funding summary, email apply@thelivingcraft.ai and mention LC-OCT-D17.	Put yourself in the manager's chair for a moment. "Can you fund this course?" is a request for money. It leaves quite a lot of work for the person reading it. Which responsibility will it help with? What do you want to practise? How much time will you need? Where could you use the learning? I would answer those questions in the request itself. For example: you review agent designs, and you want to strengthen how you judge the evidence behind a team's choices. The Living Craft includes practical building, continuous feedback and 30 live hours plus independent work. Confirm the schedule and independent-work expectation before asking for final funding approval. The manager summary puts the individual sponsorship route into one conversation. Use it to discuss employer funding for your open-cohort application. For the employer-funding summary, email apply@thelivingcraft.ai and mention LC-OCT-D17.	Perspective-taking scenario; no actual manager dialogue claimed. Keeps sponsor summary CTA, individual route and commitment. Sunil to confirm tone. Pending review. Original active.

LC-OCT-D19 One team or an industry case?

Day 19 / Carousel / Engineering sponsor / Original active

Current version	Proposed rewrite	Review note
I would separate domain relevance from company-specific consulting before discussing a team programme. Industry examples can make the learning relevant without requiring confidential company information. A company case needs a coherent team, a scoped problem and permission to use the material. The group's headcount alone cannot establish that scope. We need to understand the learning objective and what would count as useful work within the programme. Request an enterprise scoping conversation before assuming a company problem is included. For enterprise scoping, email apply@thelivingcraft.ai and mention LC-OCT-D19.	Imagine a team arriving at a learning programme with the sentence: "We have a company problem we could use." That could mean a useful example. It could also mean a production project with three other teams waiting for it. We need to know which conversation we are having. An industry case can make practice relevant without bringing confidential company material into the room. A company-specific case needs a coherent team, a clear problem and permission to use the information. The headcount cannot answer those questions for us. The learning objective and agreed scope can. Before assuming your company problem is included, request an enterprise scoping conversation. We can work out which case approach fits the team. For enterprise scoping, email apply@thelivingcraft.ai and mention LC-OCT-D19.	Hypothetical team request; three other teams is illustrative, not a delivery commitment. Preserves industry default, permissions and enterprise scoping route. Pending review. Original active.
Creative copy Choose the right case Domain relevance and company consulting have different boundaries. Industry case by default Examples and assignments reflect the domain. Keep a coherent problem A company case needs a team working on shared context.	Creative copy A learning case or a company project? Start by agreeing what the team actually needs. Make the domain familiar Industry examples and assignments are the default. Share a coherent problem A company case needs a team working from shared context.	Same format and cohort purpose.

Current version	Proposed rewrite	Review note
Agree information access Use only material the organisation permits for the programme. Scope the deliverable Headcount alone does not determine consulting scope. Discuss the team's need Enterprise programme enquiries: apply@thelivingcraft.ai	Check permission Use only information the organisation permits for this purpose. Define the work Headcount does not settle what a company-specific case includes. Discuss the case first Enterprise scoping: apply@thelivingcraft.ai	

LC-OCT-D21 Before the agent changes something

Day 21 / Video / Member / Original active

Current version	Proposed rewrite	Review note
Sending a message, granting access and issuing a refund all change something outside the model. In an illustrative review, I would examine authority, evidence, duplicate requests and uncertain tool results before allowing the action. The control needs to exist where the action happens. The Living Craft connects these questions to practical building and design critique. Explore the open cohort if you want to guide such decisions. Programme and application: https://learning.thelivingcraft.ai/?utm_source=linkedin&utm_medium=organic_social&utm_campaign=cohort_enterprise_v3&utm_content=lc-oct-d21	The tool timed out. Did the refund fail, or did only the reply fail? In this illustrative situation, pressing retry is suddenly an architecture decision. Once an agent can change something outside the model, we need to know who permits the action, how duplicates are recognised and what happens when the outcome is uncertain. We connect those questions to building and design critique at The Living Craft. Explore the open cohort if you want to guide decisions like these. Programme and application: https://learning.thelivingcraft.ai/?utm_source=linkedin&utm_medium=organic_social&utm_campaign=cohort_enterprise_v3&utm_content=lc-oct-d21	Illustrative timeout example. Preserves external-action boundaries, duplicate handling and invitation. Technical review needed before recording; no universal retry policy claimed. Pending review. Original active.
Spoken script Before an agent changes something outside the model, I want to know how that action is controlled. Think about sending a message, granting access or issuing a refund. If the tool times out, the agent may not know whether the action succeeded. Simply trying again can create a different problem. In an illustrative review, I would ask four questions. Who grants the authority? What evidence is required? How do we recognise a repeated request? And who decides what happens when the result is uncertain? Those answers should shape the tool and workflow. A reassuring sentence in a	Spoken script Imagine asking a tool to issue a refund. You wait. Then you get a timeout. Did the refund fail? Or did the refund happen and the response disappear? If we cannot tell, “try again” is not much of a plan. We might be asking the system to pay twice. This is an illustrative example, but the same question appears whenever an agent changes something outside the model: sending a message, granting access, moving money. I would start with the authority to act, the evidence required and a way to recognise the same request. Then I would make the uncertain outcome an explicit state, with a defined way to check or escalate it.	Same format and cohort purpose.

Current version	Proposed rewrite	Review note
prompt does not establish the action's actual boundary. The team needs observable behaviour it can examine and explain. At The Living Craft, we connect these decisions to practical building, evaluation and continuous feedback. The aim is to develop the ability to design and critique agentic systems and guide the team's architecture decisions. If that matches the work you want to own, explore the open cohort.	The control has to exist in the tool and workflow. We should be able to inspect what happened, not just hope the prompt covered it. At The Living Craft, members build an agentic system and connect behaviour like this to architecture reasoning, evaluation and continuous feedback. Explore the open cohort if you want to strengthen how you design, critique and guide those decisions.	

LC-OCT-D23 When self-study may be enough

Day 23 / Written / Member / Original active

Current version	Proposed rewrite	Review note
Self-study can be a good choice when you know what you need to learn and can evaluate your own work. Live work becomes useful when you want another experienced perspective on the assumptions you are making. Can you explain a decision, hear a challenge and revise the design in response? That is the experience I want to make concrete at The Living Craft: a working system, evidence and a useful architecture discussion. It should complement the reading and building you already do. Explore the programme if that kind of feedback matches what you need next. Programme and application: https://learning.thelivingcraft.ai/?utm_source=linkedin&utm_medium=organic_social&utm_campaign=cohort_enterprise_v3&utm_content=lc-oct-d23	Sometimes the best next step is to close the course tab and build the thing. If you know what you need to learn and can judge your own work, self-study may give you exactly what you need. The harder moment is when the system works and you are still unsure whether the design makes sense. Imagine explaining it to another engineer who asks one question you had not considered. Now you have a new assumption to examine, or a reason to change your approach. That is where live practice and feedback can help. At The Living Craft, I want the working system to give us something concrete to discuss: the choices, the evidence and what to try next. It should build on the reading and experimenting you already do. Explore the programme if that is the kind of feedback you are looking for. Programme and application: https://learning.thelivingcraft.ai/?utm_source=linkedin&utm_medium=organic_social&utm_campaign=cohort_enterprise_v3&utm_content=lc-oct-d23	Retains the original honest self-study alternative. Hypothetical peer question; no promised outcome or invented anecdote. Sunil to confirm comfort with the direct opening. Pending review. Original active.

LC-OCT-D25 A useful application explains the starting point

Day 25 / Carousel / Member / Original active

Current version	Proposed rewrite	Review note
An application should help us understand where you are starting and what you want to develop. Describe your system-design experience, the responsibilities you hold and one question you want to work on. Explain whether you are self-funded or need employer approval. There is no need to attach confidential company architecture. Fit, schedule and commercial terms should be clear before anyone commits. Apply to discuss your open-cohort fit. Programme and application: https://learning.thelivingcraft.ai/?utm_source=linkedin&utm_medium=organic_social&utm_campaign=cohort_enterprise_v3&utm_content=lc-oct-d25	“I want to learn AI” is a starting point. It leaves us with a few questions. “Here is the kind of system I work on, and here is a design decision I want to get better at” gives us a conversation. That is what a useful application should do. Help us understand your starting point. Tell us about your system-design experience, the responsibility you hold and one question you would like to work on. Let us know whether you are self-funded or need employer approval. There is no need to send confidential company architecture. An explanation in your own words is enough to begin. We can then discuss fit, schedule and terms before you commit. Apply to explore your place in The Living Craft’s open cohort. Programme and application: https://learning.thelivingcraft.ai/?utm_source=linkedin&utm_medium=organic_social&utm_campaign=cohort_enterprise_v3&utm_content=lc-oct-d25	Generic example application language; not a quote from an actual applicant. Preserves prerequisites, funding and application CTA; six panels. Pending review. Original active.
Creative copy Give your application context A clear starting point supports a useful fit conversation. Describe your experience Which systems and design decisions have you worked with? Name the responsibility What do you need to design, critique or guide?	Creative copy Give us a starting point A useful application opens a conversation. Describe your experience Which systems and design decisions have you worked with? Name the responsibility What do you need to design, critique or guide?	Same format and cohort purpose.

Current version	Proposed rewrite	Review note
Explain your learning goal Where would practice and feedback help? Clarify the buying route Self-funded or employer-funded. Keep confidential materials out. Apply for the open cohort Start the fit conversation: learning.thelivingcraft.ai	Bring one question Where would practice and feedback help you? Explain the funding route Self-funded, employer-funded or still deciding. No confidential files needed. Start the conversation Open-cohort applications: learning.thelivingcraft.ai	

LC-OCT-D27 Review the team's next design

Day 27 / Video / Member / Original active

Current version	Proposed rewrite	Review note
Guiding a team's design starts with understanding the decision it is trying to make. I would ask the team to explain the task, the boundaries and the evidence behind a choice. Then we can locate the uncertainty and identify the next useful change. The Living Craft connects this review practice with a working agentic system and continuous feedback. Explore the open cohort to develop how you design, critique and guide. Programme and application: https://learning.thelivingcraft.ai/?utm_source=linkedin&utm_medium=organic_social&utm_campaign=cohort_enterprise_v3&utm_content=lc-oct-d27	Imagine a design review where everyone leaves with more comments and nobody knows what to do next. The document got longer. The decision did not get easier. I would bring the conversation back to the task, the boundaries and the evidence. Where is the uncertainty? What is the next useful thing to check? Guiding a team means helping it answer those questions. At The Living Craft, we connect that practice to a working build and continuous feedback. Explore the open cohort to develop how you design, critique and guide. Programme and application: https://learning.thelivingcraft.ai/?utm_source=linkedin&utm_medium=organic_social&utm_campaign=cohort_enterprise_v3&utm_content=lc-oct-d27	Illustrative design meeting, not a claimed event. Keeps reviewer guidance, practical work, feedback and commitment. Sunil to validate his preferred phrasing. Pending review. Original active.
Spoken script When a team brings me a design, I first want to understand the decision it is trying to make. What task should the system perform? Which constraints matter? Where can it act, and where should it stop? Then I ask the team to explain the evidence behind its choices. The review becomes useful when we can locate an uncertainty and agree the next piece of work. Perhaps we need to inspect a failed tool action, narrow a permission or compare two approaches against the same task.	Spoken script Imagine finishing a design review with a document full of comments. Everyone contributed. Then someone asks, "So, what do we do next?" If we cannot answer that, we still have work to do. I would start by asking which decision the team is trying to make. What task should the system perform? What constraints matter? Where can it act, and where should it stop? Then we can look at the evidence behind the proposed approach. Perhaps a tool failure needs closer inspection. Perhaps a permission is too broad. Perhaps we	Same format and cohort purpose.

Current version	Proposed rewrite	Review note
Guiding the team means helping it make that decision clearly. It does not require replacing every choice with my preferred design. At The Living Craft, members build a working agentic system and connect the work to architecture reasoning, evaluation and system-design review. Continuous feedback identifies strengths, gaps and follow-up work. The programme includes 30 live hours plus independent work. Explore the open cohort if you want to strengthen how you design, critique and guide a team's architecture decisions.	should compare two approaches on the same task. The next step should be concrete enough for the team to take and bring back. That is more useful than replacing every choice with the reviewer's favourite design. This is an illustrative review situation. At The Living Craft, members build a working agentic system and practise connecting it to architecture reasoning, evaluation and design review. Continuous feedback helps identify strengths, gaps and follow-up work. The programme includes 30 live hours plus independent work. Explore the open cohort if you want to strengthen how you design, critique and guide a team's architecture decisions.	

LC-OCT-D29 Bring a question you can examine

Day 29 / Written / Member / Original active

Current version	Proposed rewrite	Review note
A useful starting question is concrete enough to investigate. For example: what must be true before this agent can change external state? What evidence would make me revise the workflow? How would I explain the design to another engineer? You can start from an industry case; you do not need to disclose a company's confidential system. The programme connects a working build with architecture reasoning, evaluation and review. If you are considering joining, describe your experience and the question you want to work on in the application. We can then discuss fit and the full commitment. Programme and application: https://learning.thelivingcraft.ai/?utm_source=linkedin&utm_medium=organic_social&utm_campaign=cohort_enterprise_v3&utm_content=lc-oct-d29	A small, awkward question can give you more to work with than a very ambitious project title. For example: what happens if this agent is asked to take the same action twice? You can trace the request, inspect the tool behaviour and decide what must change. The question gives the work a shape. Or start with something else you genuinely want to understand: what evidence would make you revise the workflow? How would you explain a design choice to another engineer? You can explore those questions through an industry case. There is no need to disclose a confidential company system. The Living Craft connects that practical build to architecture reasoning, evaluation and review. If you are considering joining, describe your experience and the question you want to examine in the application. We can then discuss fit and the full commitment. Programme and application: https://learning.thelivingcraft.ai/?utm_source=linkedin&utm_medium=organic_social&utm_campaign=cohort_enterprise_v3&utm_content=lc-oct-d29	Illustrative investigation prompts, no personal incident. Preserves the existing concrete-question and application intent. Sunil to confirm the closing sounds natural. Pending review. Original active.