

The Living Craft

Website growth and search roadmap

STATUS	OWNER	LAST UPDATED
Prepared for review	Team Alchemy	10 September 2026

Authors	Team Alchemy
Reviewers	Sunil Mathew and Ein
Related docs	Cohort handoff package
Scope	Page architecture content clusters and search readiness

Prepared by Team Alchemy for Sunil and Ein. 10 September 2026. The immediate delivery is the cohort landing page and pipeline. This roadmap gives the wider learning and architecture-advisory website a coherent structure without changing the current campaign or delaying the cohort release.

PAGE HIERARCHY AND NAVIGATION

Proposed paths below are targets for inventory and design, not confirmed live URLs. Keep learning.thelivingcraft.ai working until routing and redirects are approved and tested. Inventory the existing main domain, field notes, assessment and CAIO surfaces before making any migration decision; do not silently delete or promote them.

Home / → About /about/; Programmes /programmes/; Advisory /advisory/; Resources /resources/; Tools /tools/; Contact /contact/.

Programmes → Agentic systems cohort /programmes/agentic-systems/; Enterprise learning /programmes/enterprise/.

Resources → Guides /resources/guides/; Field notes /resources/field-notes/; Templates /resources/templates/.

Tools → Agent Design Check /tools/agent-design-check/; Architecture Review Skill /tools/architecture-review-skill/.

Utility routes → Apply /apply/; submission confirmation /application-received/; communication preferences /communication-preferences/; privacy /privacy/; terms /terms/.

Header order: Programmes, Advisory, Resources, Tools, About, then Explore the cohort. Logo returns home. Footer groups programmes, resources, company/contact and privacy/terms. Guide pages use breadcrumbs and contextual related links. Every published page has at least one inbound internal link. Do not publish empty hubs merely to fill the navigation.

PAGE BRIEFS BY INTENT AND CONTEXT

Each row includes a sentiment hypothesis, not a diagnosis of individual visitors. Replace hypotheses with observed questions from enquiries and conversations over time.

Page	Reader intent and context	Concern hypothesis and answer	Evidence and next step
Home	Engineer or sponsor discovering the practice	Wants to understand what is available; explain learning, resources and advisory clearly	Verified Sunil background, practical example; explore cohort or resources
About	Reader checking who is teaching/advising	Sceptical of vague authority; connect verified experience to current approach	Source-backed roles and approved current work; programme or advisory
Programmes	Visitor choosing individual or group learning	Unsure which route fits; explain cohort, individual sponsorship	Current offer facts; corresponding programme page

Page	Reader intent and context	Concern hypothesis and answer	Evidence and next step
		and enterprise separately	
Cohort	Experienced practitioner assessing relevance	Wary of another generic course; show work, feedback, prerequisites and limitations	Validated curriculum, VSL, example; application or enquiry
Enterprise	Sponsor/L&D exploring a group	Needs domain relevance and practical scope; explain readiness, group learning and coordination	Approved scope and industry-default facts; enterprise enquiry
Advisory	Leader facing a design decision	Wants useful guidance without unclear project scope; explain framing, review and next-step planning	Sunil-approved service boundaries; scoping conversation
Resources	Practitioner trying to solve a design question	Overwhelmed by tool lists; organise by problem and task	Authored guides and worked examples; relevant tool or deeper guide
Guides	Reader needs a comprehensive explanation	Wants an answer they can apply; direct response, context, example, evidence, limitations	Named author/reviewer and primary sources; related guide/tool
Field notes	Practitioner following current building lessons	Wants specifics and honest uncertainty	Verified observation or labelled illustration; discussion or guide
Templates	Engineer needs to run a review or request funding	Wants something usable immediately	Completed example and blank usable artifact; browser checker or programme
Tools	Practitioner wants to try the method	Wants value without a sales form	Transparent rubric and sample output; optional

Page	Reader intent and context	Concern hypothesis and answer	Evidence and next step
			resources/cohort/advisory
Apply	Candidate ready for a fit discussion	Concerned about commitment or confidential material	Clear fields, privacy and process; saved application confirmation
Contact	Visitor has a specific unresolved question	Needs the right person without repeating context	Enquiry receipt and route ownership; human response

Advisory wording is limited to use-case framing, architecture review and next-step guidance for this roadmap. Do not add implementation guarantees, fixed delivery packages or Fractional CAIO promotion without separate scoping. Sunil confirms service facts before publication.

FOUR CONTENT CLUSTERS

Publish one pillar guide and three supporting articles per cluster. These sixteen pieces are editorial briefs, not claims of existing articles. Proposed slugs use /resources/guides/ for pillars and /resources/field-notes/ for supporting material.

Cluster and pillar slug	Supporting article titles and slugs	Useful artifact
Designing agentic systems: agentic-system-design	When a workflow is enough: workflow-or-agent; Who may call the tool: tool-permissions; What to do with uncertain evidence: uncertain-evidence	Agent-design canvas and browser checker
Evaluating agent behaviour: agent-evaluation	Turn failures into test cases: failures-to-test-cases; Evaluate retrieval and answers separately: retrieval-and-answer-evaluation; When an automated judge disagrees: reviewing-evaluators	Evaluation-planning workbook
Reliability security and cost: agent-reliability	Did the action happen twice: retries-and-duplicates; Plan intervention and recovery: intervention-and-recovery; The cost around the model: workflow-cost	Failure-review checklist
Guiding an engineering team:	Run a useful architecture review: useful-design-reviews; Record the	Design-review agenda and decision record

Cluster and pillar slug	Supporting article titles and slugs	Useful artifact
agent-design-review	trade-off: architecture-decision-record; Choose the first use case: choosing-an-agent-use-case	

Each pillar answers the broad question and links to its three supporting pieces. Each supporting piece links back to the pillar and to relevant neighbouring material. Link to a tool where it helps the task; add programme/advisory links when there is a genuine next step. Maintain one canonical page for a question rather than cloning copy for small keyword variations.

First resource release: agent-design pillar, its three supporting articles, review workbook and browser checker. Later releases complete evaluation, reliability and team-review clusters in that order. Expand an ebook only when enough original work exists to warrant one.

AGENT DESIGN CHECK FUNCTIONAL SCOPE

Offer an ungated browser checklist. Use structured yes/no/not-yet-defined answers about purpose, evidence, tool permissions, external actions, evaluation, uncertainty/recovery and ownership. Include a built-in illustrative example so visitors can understand each question without uploading company materials.

The result shows which questions remain unanswered, why they matter and a prioritised set of next steps with guide links. Distinguish a reported gap from a missing answer. Use transparent rules reviewed by Sunil, not a model-generated readiness score. All-unknown input returns a starting checklist, not a failure grade. A claim that every box is addressed still requests evidence rather than certifying production readiness.

Process answers in the browser. Offer a downloadable text/print summary without login; retain no architecture answers on the server by default. Optional email signup follows the result and is separate from using the tool. Track starts/completions only with appropriate analytics permission. No lead record is created from anonymous use.

After rubric approval, package the same method as Claude and Codex architecture-review skills. Review only user-selected material; return referenced observations, questions, limitations and next actions. Include a sample input/report, usage guidance and versioned rubric. Verify each host independently; no cross-host compatibility claim before tests. Do not execute or modify the user's code as part of the review by default.

SEARCH AND ANSWER QUALITY REQUIREMENTS

For each page, author the direct answer first, then who/when it applies, a concrete explanation or worked example, supporting evidence, limitations and the appropriate next step. Put important text in accessible HTML. Video transcripts, tables and diagrams support the explanation rather than replacing it. Publish author/reviewer identity and meaningful revision dates; never refresh dates without content review.

Use descriptive page titles, one clear primary heading, semantic sections, descriptive links, canonical URLs, a current XML sitemap and crawlable internal links. Validate mobile readability and content parity between rendered HTML and source. Structured data may describe the organisation, person, article, breadcrumbs, course or video only when the visible content supports it. Do not manufacture reviews or expect FAQ markup to secure rich results.

Google states that its AI features use the same SEO foundations and require no special AI files or markup. Search snippets must remain eligible; robots, CDN rules and rendering can block access. OAI-SearchBot and GPTBot serve separate search/training purposes; Perplexity also documents its search crawler. Verify search access using the current official guidance without changing training preferences by implication. These measures support discoverability, not a guarantee of rankings or citations.

Sources checked during this task's planning on 10 September 2026: Google AI features (<https://developers.google.com/search/docs/appearance/ai-features>), Google helpful content (<https://developers.google.com/search/docs/fundamentals/creating-helpful-content>), Google video guidance (<https://developers.google.com/search/docs/appearance/video>), OpenAI crawlers (<https://developers.openai.com/api/docs/bots>), Perplexity crawlers (<https://docs.perplexity.ai/docs/resources/perplexity-crawlers>).

The resource/skills pattern has a relevant example in Hamel Husain and Shreya Shankar's evaluation skills (<https://hamel.dev/blog/posts/evals-skills/index.html>). It supports the format choice; The Living Craft must publish its own reviewed method. Justin Welsh's founder-story guidance (<https://justinwelsh.me/essays/build-a-better-founder-story>) informs narrative structure using true experiences, not fabricated biography.

TWELVE QUESTION VISIBILITY SAMPLE

Use the same wording at baseline and monthly review. These are selected research questions, not validated search-volume targets. Log platform, date, region, language, model/mode where visible, exact response evidence, brand mention, cited URL, accuracy and competitor mentions. If an engine is unavailable, leave it untested. Keep organic referrals, AI referrals and manual citation observations separate. Google's AI-feature traffic is included in its overall Web search reporting; do not invent a clean AI Overview traffic split.

1. When should an engineering team use an agent instead of a fixed workflow?
2. How should an agent's tool permissions be designed?
3. What should an agent do when its evidence is uncertain?
4. How do you turn agent failures into evaluation cases?
5. How do you evaluate retrieval separately from an agent's final answer?
6. How do you check whether an automated evaluator is useful?
7. How can an agent avoid repeating an external action after a timeout?
8. What should an agent recovery and human-intervention plan include?
9. How do you estimate the full cost of an agentic workflow?
10. What questions belong in an agentic architecture review?
11. How can an employer evaluate funding for agentic-system learning?
12. What does The Living Craft offer experienced engineers and engineering teams?

ROADMAP ACCEPTANCE AND ASSUMPTIONS

Before a resource is published, confirm its example, sources, author, answer quality, links, mobile rendering and appropriate next step. Checker acceptance covers unknown/conflicting input, repeatability, transparent rules, exports and no unsolicited data collection. Skill acceptance covers evidence references, missing material and host compatibility.

The wider site remains a phased roadmap; this handoff does not claim sixteen articles, a browser checker or installable skills have been built. Ein's immediate implementation remains the cohort page and pipeline. Existing URLs, traffic and infrastructure could not be verified live during planning, so no migration or visibility baseline is asserted.